

Dwelling Age Manual Review (DAMR)

The end-to-end briefing: what the alert is, what we tried, what the census found, and the decisions it opens up

Executive summary

DWELLING_AGE_MANUAL_REVIEW is Steadily's highest-volume underwriting alert. Carrier guidelines require a pre-bind review of every dwelling 101+ years old, so the alert blocks every flagged quote until an agent refers it to underwriting and UW responds. A July effort to reduce the alert inside the SSB sprint tried three fixes and shelved all three; the work pivoted to a full quantitative census of what the review actually does (BUC-4962, Will Henry). That census is now complete: every eligible quote since Jan 15 traced end to end, every field change attributed to an actor, and **all 33,696 underwriting review notes classified** and cross-checked against the audit log. The headline: **roughly 2% of reviews find an age-driven problem**. The gate's dominant effect is attrition — 55% of eligible flagged business never reaches UW, and never-referred business never binds. When asked, UW approves 83% of the time (91% of decided outcomes), a third of the time with nothing recorded, half the time changing nothing on the quote. The rejections that justify the gate are real, but ~82% of decline substance is not about age — it is commercial exposure, historic districts, and unit counts, each detectable from data signals of its own. The roof-surfacing exclusion, the review's most common substantive act, is already applied as a de facto default. These findings frame a concrete set of decisions — targeted decline rules, formalizing the roof exclusion, replacing the note-writing referral funnel with a structured application — now in front of underwriting, product, and carriers.

1 · What the alert is and why it exists

DWELLING_AGE_MANUAL_REVIEW fires on any dwelling 101 years or older. It exists because carrier guidelines require pre-bind review of these risks: old homes carry elevated odds of knob-and-tube wiring, aged roofs, structural issues, historic designations, and stale public records that make automated dwelling assumptions unreliable. It is the single highest-volume underwriting alert — it appeared on 67,053 new-business quotes in the six-month study window alone — and it blocks binding on 100% of flagged quotes until an underwriter responds to a referral.

One mechanic matters more than any other: **the review is not a yes/no gate**. When UW approves, it often modifies the quote first — applies a roof-surfacing exclusion, corrects square footage or property type, adjusts coverages — with the fields frozen so agents can't undo the change. Those edits carry rate and coverage impact. Skip them, and mispriced policies get reformed at inspection or re-rated at renewal. Every simple “just let them through” idea has died on this mechanic.

2 · How we got here (July 2026)

The SSB sprint (Jul 4–14, run by Mitch Thomson) took the first swing, prompted by an analysis from Julie Scully: of SSB quotes blocked *solely* by this alert (May 1–Jun 30, n=206), only 38% were ever submitted to UW — but **90% of submitted quotes were approved**, while 62% simply dropped off. If nearly everyone passes, why block anyone? Three fixes were explored:

Option explored	What it was	Why it was shelved
Remove the alert	Let 101+ homes bind without pre-bind review.	A 30-quote hand audit of never-referred quotes (Jenna Goldband, UW Director) found only 50% clean approvals; 23% would be approved with changes , 23% needed more information, 3% rejected. The 90% headline ignored conditional approvals.
CX chat assist	Live chat widget on the Coverages page; CX resolves the alert in ~10 minutes.	Judged too disruptive to the sales flow. Tickets closed Won't Fix.
Auto-refer + email link back	Auto-create a referral for single-alert quotes; email customers a link when approved clean.	If UW modifies the quote, coverage and premium change with no agent present to explain it. Expected upside was only ~1–2 binds/week.

The sprint closed Jul 14 with all three marked “won't do,” and a carried-forward mandate from LaNae Rueda (CXO, carrier guidelines owner): solve this for SSB *and* the rest of the MGA — it simply needed more lift than one sprint week. Her tradeoff frame: fewer blocks (lower OpEx) vs. more binds (higher GWP) vs. more post-bind UW action (bad for retention).

At the Jul 21 Product/UW sync the approach changed: **stop trying to eliminate the alert; understand and automate what the review actually does.** LaNae's motivating example — if the most common approval action is applying a roof-surfacing exclusion, auto-apply it on every 101+ home and let agents refer only to *remove* it with evidence. Will Henry took ownership of the analysis (BUC-4962), scoped NAUI-wide rather than SSB-only, and built on real referred quotes with real UW responses rather than counterfactual hand-reviews.

A note on “paired alerts”

Every version of this work restricts to quotes where the age alert is the *only* underwriting question. If a quote carries another review-requiring alert, it goes to UW regardless, so fixing this alert changes nothing for it. In the final census, 58% of flagged quotes are age-alert-only; the 42% that co-fire (usually with exactly one other alert: another roof question, owner-occupancy, or a Coverage A threshold) would see their review *shrink* rather than disappear.

3 · The census: how the analysis was built

Everything below comes from a frozen, reconciled census — not a sample — published Aug 4, 2026 (“Dwelling-Age Pre-Bind Review — Findings”, Will Henry) with receipts in Metabase collection 6787 (questions 01–16, validations A1–A10). In plain language:

- **Cohort.** All sMGA new-business quotes carrying the alert, originated Jan 15 – Jul 22, 2026: 67,053 quotes. Excluded with separate receipts: carryover/rewrites (5,470), self-serve channel (~1,860 — no agent exists to refer), immature quotes (~11,100 born within 35 days of the freeze — the 95th-percentile full lifecycle), unattributable (2). Result: 48,589 quotes, of which **28,121 are age-alert-only** — the clean test of what this alert alone does.
- **Outcomes.** A review's outcome is its last formal approve/reject verdict; post-verdict follow-up chatter doesn't overwrite it (that convention had been burying 124 approvals and 56 rejections). “Stalled” means no verdict ever.

- **Change detection.** Field-level diffs over every tracked property and coverage column, per quote version, from first referral to final decision — each change attributed to underwriter, agent, or staff via the audit log. Unknown actors: 0.04% of changed quotes.
- **The note census.** All 33,696 underwriting review notes on referred quotes, classified note-by-note (with thread context) against a frozen codebook, with strict structured outputs and zero errors; 99.3% accuracy on a hidden disposition canary; label-vs-audit-log coherence receipted. Coverage: 12,603 of 12,605 referred quotes.
- **Gate integrity.** No organic business issues while the alert is live — the handful of exceptions are internally created clone quotes issued via the explicit UW override workflow. “Never referred” therefore means “never bound.”
- **Robustness.** The headline unchanged-approval rate moved within a one-point band (45.4% → 46.5%) across six successively tighter cohort definitions. Data Science independently reproduced the approval-rate reconciliation.

4 · What the census found

4a · The funnel: most of what the gate touches dies without ever being judged

Stage (age-alert-only cohort)	Count	Share
Age-alert-only quotes (sMGA NB, Jan 15 – Jul 22)	28,121	—
Never referred to UW	15,447	54.9%
Agent referred to UW	12,605	44.8%
Approved (last verdict = yes)	10,430	82.7% of asked
Rejected	1,059	8.4% of asked
Stalled — no verdict, ever	1,116	8.9% of asked
Issued, all paths (final status)	5,196	18.5%

- **Reconciling the 90%.** Count only decided outcomes and the approval rate is 90.8–91.1% (independently reproduced by Data Science). The earlier 90% wasn't wrong — it discarded the stalled tenth, which consumes underwriter attention and produces nothing: 99% of stalled quotes sit untouched for 35+ days, median ~15 weeks.
- **The 55% drop-off is the referral mechanism, not this alert.** Benchmarked against every other alert, note-writing referral gates lose half or more of everything they touch; review-requiring alerts run 15–42% referral rates and this alert sits at the top of that band at 42% — on 4x the volume of the next-largest alert. Agents engage it more than almost anything, plausibly because it's known to be winnable.
- **The precedent that this is fixable.** The water-claims alert ran a 40.6% referral rate under the same note-writing flow. Replaced with a structured application, ~85% of flagged quotes now clear that alert, versus ~40% that ever reached UW before. Same book, same risk — the funnel, not the risk, was doing the blocking.

4b · The yes is hollow: half of approvals change nothing on the quote

Outcome of the 10,428 analyzable approvals	Quotes	Share
No change of any kind (literally zero field diffs)	4,848	46.5%
Roof-surfacing exclusion only	1,463	14.0%
Roof exclusion + roof detail corrections	119	1.1%
Roof exclusion + property corrections	569	5.5%
Roof exclusion + coverage changes	488	4.7%
Changes without roof exclusion	2,938	28.2%
Limit recompute only / attribution unclear	5	0.05%

- **1 in 4 approvals ends in roof-surfacing treatment** (2,639 quotes) — the modal substantive output of the entire review process is one standardized coverage action.
- **~85% of roof-exclusion approvals show no roof fact changing** between referral and exclusion. The exclusion is applied by default, not discovered. ~72% of approvals resolve in under one hour; the modal review is: agent refers, UW approves within minutes, nothing changes.
- **Where UW does engage, it is real work.** UW is the dominant actor on property-fact corrections (roof type, square footage, construction); agent-attributed changes concentrate in coverage tuning that happens during, not because of, review.

4c · Even the paper trail treats it as a formality

Every approval note was classified for what the yes rested on:

Basis for the yes	Approvals	Share
UW did work — verified externally, corrected data, carved out risk	4,609	44.2%
Relied on agent evidence (documents, attestation, signed form)	1,202	11.5%
Deferred — post-bind inspection or prior precedent	1,016	9.7%
Nothing recorded — rubber stamp or empty note	3,603	34.5%

The 1-in-4 finding

Cross-examined against an independent witness — the field-level audit log: where notes claim work, the log shows UW-made changes 92% of the time; where notes record nothing, the log shows no UW-attributed change 89.6% of the time. Net: **3,230 approvals — 31% of approvals, 26% of everything agents refer — have no recorded basis AND no underwriter-attributed change to the quote.** 71% of approvals arrive on UW's very first touch. And agent evidence is ceremonial to the verdict: approval runs 82–85% at every evidence tier (nothing, attestation, documents); evidence only reduces stalls (9.5% → 6%). UW decides from its own desktop research.

4d · The no is real — but it isn't about age

The 1,059 rejections are the gate's entire measurable defense. Every note carrying decline substance was read — 1,448 notes, including declines delivered inside back-and-forth threads that never became formal rejections

(12% of decline substance travels that way):

Decline reason	Share of substance	Age-related	Offers cure path
Commercial exposure on the property	28.2%	0%	19%
Historic district	15.4%	0%	4%
Too many units	12.8%	0%	7%
Ineligible construction material	7.3%	95%	25%
Condition seen in aerial/street imagery	6.4%	99%	66%
Wrong product for renovation state	6.4%	1%	82%
Business operations ineligible	4.4%	0%	19%
Pre-1940 systems / electrical	4.2%	100%	16%
Twelve smaller categories	15.0%	mixed	mixed

~82% of decline substance is not age-driven. The top reasons — commercial exposure, historic districts, unit counts — are real risks the age tripwire happened to surface, each decidable from data signals of its own (property-use records, historic-registry geography, parcel and unit data). The genuinely age-driven declines (~18% of substance, roughly **1.5% of all referrals**) concentrate in construction material and visible condition — categories that usually come with a fix-and-return path rather than a firm no. Hence the report's headline: about 2% of reviews find an age-driven problem.

4e · The stalled tenth dies on unanswered photo requests

- UW's asks are reasonable — confirm occupancy and unit count (22% of asks), exterior/roof photos (21%), official property records (18%). Only 2.3% of information requests contain no codable question.
- The answers are the leak: photo requests get anything back only 45% of the time; property records 44%; ownership documents 19%. When a photo request is the last word on a thread, the thread dies 57% of the time. **Exterior/roof photos alone account for 320 dead threads** — roughly a third of the stalled tenth. 50 referrals died at the automated-requirements bot without a human underwriter ever touching the thread.

4f · Even a yes doesn't save the business

Review outcome	Quotes	Issued	% issued
Never referred	15,447	54	0.3%
Approved — no changes	4,848	2,449	50.5%
Approved — other changes	2,938	1,432	48.7%
Approved — roof exclusion only	1,463	670	45.8%
Stalled (no verdict, ever)	1,062	16	1.5%
Rejected	1,059	1	0.1%

- Half of approvals never bind — and review latency doesn't explain it: the walk rate is *highest* where review was fastest (54% for sub-hour approvals vs. 28% for week-long episodes). What UW did to the quote doesn't move issue rates either.

A possible adverse-selection effect

At identical review speed, independent-agent business walks at 55% vs. 40% for direct — a ~15-point gap exactly where an agent *could* move the customer to a carrier without the pre-bind requirement. One reading: the gate sheds placeable business and retains the risks with fewer alternatives. Caveat: the two channels write different books, and this comparison is not yet risk-adjusted.

5 · Synthesis: an attrition gate with an age-shaped trigger

Put together, the census describes a review that is not functioning as an age review:

- **Its dominant effect is attrition, not underwriting.** 55% of eligible flagged business never reaches UW and essentially never binds; another tenth of referrals stalls to death, mostly on unanswered photo requests. The gate's biggest output is deleted business no underwriter ever judged.
- **When asked, the answer is yes** — 83% (91% of decided), a third of the time with nothing recorded, half the time changing nothing, usually within the hour.
- **The review isn't consuming agent input.** Evidence doesn't move the verdict, and UW's own document requests mostly go unanswered. UW decides from its own desktop research.
- **Its defense — the 8.4% rejections — is mostly about other risks:** commercial exposure, historic districts, unit counts. Each has its own data signal. The truly age-driven declines are ~1.5% of referrals, and skew toward fix-and-return conditions.
- **Its modal substantive act is one standardized coverage action** — the roof-surfacing exclusion, applied ~85% of the time without any roof fact changing first.

6 · What this makes possible

The census does not prescribe a replacement — those are underwriting, product, and carrier decisions. What it does is map where the review's work actually happens, which frames the candidate moves. Ordered by leverage-to-lift:

- **1. Size the rule-able rejections.** How many of the 1,059 actual rejections would existing structured signals have caught — property-use/commercial indicators, unit counts, historic-district geography? The analysis flags this as the highest-value next piece of work: it is answerable from data already on or around the quote, and it directly feeds targeted decline rules that let the gate's defense survive without the gate.
- **2. Stand up the historic-registry check (DISCO-690).** Historic district is 15.4% of decline substance and a major information-request theme; a quote-time registry check mirrors the existing prohibited-business dataset pattern. The ticket sits unassigned in Discovery.
- **3. Test whether outcomes are predictable ex ante.** Do features available at quote time (roof age/condition score, year-built band, data completeness, property type, state) separate no-change vs. roof-exclusion-only vs. substantive-change approvals? If yes, a triage router — auto-clear the predictable head, refer the tail — is viable; if no, only item-level checklist automation survives.

- **4. Decide the roof-surfacing exclusion's status.** It is already a de facto default. Formalizing it (as a default or bind condition, with a remove-with-evidence path) is a judgment call the data can inform but not make — the decision needs the over-application cost sized: UW currently applies it to ~25% of approvals, so a blanket default touches the other ~75%.
- **5. Replace the note-writing funnel with a structured application.** The single biggest lever on the 55% never-referred and the stalled tenth: collect occupancy, unit count, and roof photos up front — the exact asks that currently kill threads. The water-claims alert is the working precedent.
- **6. Expand the post-bind inspection lane.** UW already defers ~10% of approvals to post-bind verification; whether that lane can carry condition checks for more of the book is an underwriting capacity question.
- **7. Complete the validation battery.** Stability probe plus a 100-thread human holdout — inexpensive, and a prerequisite before any note-census figure is quoted to carriers or used in a formal decision.

Who decides what

Owner	The decision in front of them
LaNae Rueda (CXO)	Whether the roof exclusion becomes a formal default or bind condition; whether the validated census supports a carrier conversation about relaxing or condensing the pre-bind requirement; confirmation that the note taxonomy matches underwriting reality.
Jenna Goldband (UW Director)	The 100-thread human holdout; whether the effective checklist (commercial, historic, occupancy, condition) is correctly enumerated; post-bind inspection lane capacity.
David Curry	Whether Pluribus picks up the historic-registry scrubber (DISCO-690); the census strengthened the shared-benefit case, but bandwidth is the constraint.
Rob Soule	Sequencing: which candidate moves get staffed first, and whether the structured-application path gets scoped.
Will Henry	Runs the validation battery; owns the census tables and the decline-signal sizing if greenlit.

7 · Caveats and open items

- **Pending validation.** The note-census sections carry a pending-validation badge until the stability probe and 100-thread human holdout complete. Treat those figures as directionally solid, not yet decision-grade.
- **Not risk-adjusted.** Channel comparisons (independent vs. direct walk rates, friction) do not yet control for differences in the books the two channels write.
- **Measurement notes.** Quote status is read live and drifts daily; the ~85% roof-fact and ~72% sub-hour figures were measured on an earlier, wider cohort cut (directions verified stable across every recut); the referral-rate baseline uses lighter cohort hygiene applied identically to every alert. Underwriters changed a quote without writing anything down in 10.4% of nothing-recorded approvals — measured, not assumed.
- **Open analytical items.** An abandonment baseline for the never-referred population (comparable unflagged quotes); outcomes when document asks go unanswered; quantifying how much decline signal already exists in structured quote data; TECHSUP-1802 (previously approved pre-1920 risks re-triggering at renewal — possibly fixed, unverified).

8 · People, artifacts, and sources

Who / what	Role
Will Henry	Built and ran the census (BUC-4962): cohort spine, field-level change detection, note classification via the Claude batch API; author of the Aug 4 findings report.
LaNae Rueda (CXO)	Owns carrier guidelines and the decision guide; source of the auto-exclusion concept; gate for the note taxonomy.
Jenna Goldband (UW Director)	Ran the original 30-quote hand audit whose 50% clean-approval rate independently matches the census's 46.5% no-change finding.
Mitch Thomson (PM) / Julie Scully	Ran the SSB chapter and the original analysis whose 90% figure the census reconciles.
Rob Soule / David Curry	Rob: sequencing, provisioned batch-API access. David: checklist framing and the Pluribus / historic-registry overlap.
Findings report	Claude artifact "Dwelling-Age Pre-Bind Review — Findings (Aug 2026)" — every claim chips to a numbered Metabase question.
Data	Metabase collection 6787 "Dwelling Age Manual Review": questions 01–16 + validations A1–A10. Tables (dbt_dev): damr_census_flat_20260803, damr_note_classification_20260731, damr_note_docs_20260803, damr_note_meta_20260803, damr_eligibility_map_20260728, damr_channel_map_20260724.
Tickets	BUC-4962 (analysis) · DISCO-690 (historic registry — Discovery, unassigned) · TECHSUP-1802 (renewal re-trigger) · DISCO-702 / GLD-1511 (Won't Fix, SSB era).