Dwelling Age Manual Review (DAMR)

The pre-fill flip: the test, what it validates, and what each outcome means · v8 · prepared Aug 7, 2026

TL;DR

The proposal (from the Suph / Christine / Cameron discussion): on 101+ dwellings, stop pre-filling the property attributes UW routinely corrects and make the agent enter them, paired with the conditional roof-surfacing exclusion (RSE) default already adopted. During the test **everything else stays the same** - the alert fires, blocks, the agent refers, UW reviews every referral. Only the data's provenance changes.

The hypothesis it tests: **UW's modification work is caused by our bad pre-fill, not by underwriting judgment.** The prediction is specific - untouched approvals climb from today's ~45-50% toward a measured ceiling of **89%**. The ceiling is 89% and not 100% because ~11% of approvals carry deliberate coverage work that better data cannot remove.

If it succeeds: the review's data job is gone, the carrier case for the approve path is made, and county records demote to a backstop. If it fails: the agent cannot be the source of record, and county/assessor matching (Rob's build-vs-buy) becomes the critical path. Either way, the rejection leg is untouched: the alert cannot die until the Pluribus/Spotlight decline signals ship with receipts (commercial in build; historic 6-12 weeks out).

Preface · The frame: the problem, the goal, and three locks

The problem. Every 101+ dwelling is blocked behind a human - about 5,200 quotes a month where dwelling age is the only question on the quote. The review the block buys mostly says yes, slowly, and touches little. Meanwhile 60% of the blocked business dies without ever being reviewed. It is an expensive gate that mostly says yes, and its biggest output is attrition.

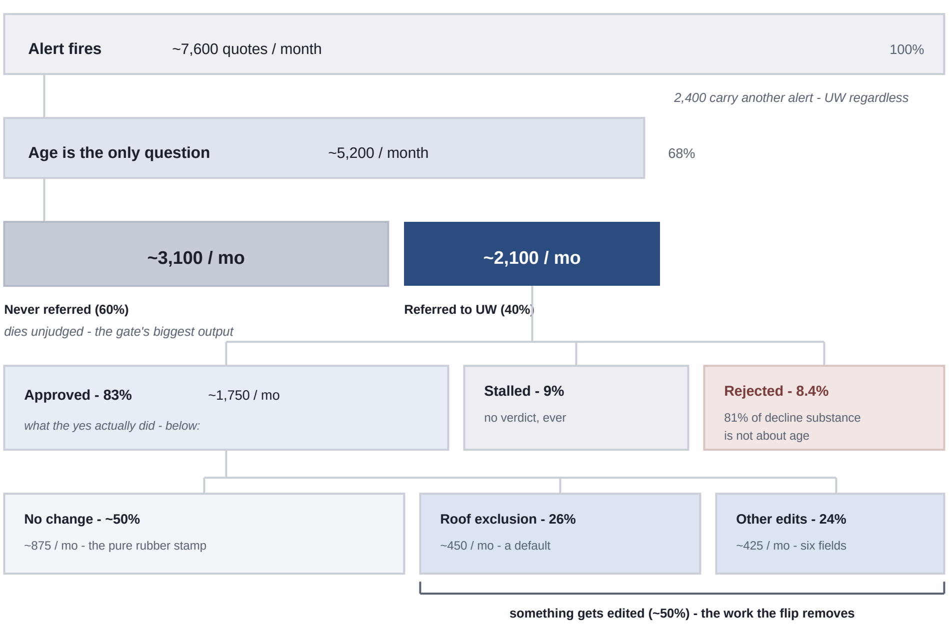

The funnel, monthly: the block deletes more business than it judges - and when asked, the answer is almost always yes.
The bottom row is the test's target: RSE defaulted + agents on the six fields = 89% of approvals need zero UW edits.

The goal. Stop blocking the quotes that end in yes - recover the dead 60% (~3,100 quotes a month), free the underwriters - while still catching the 8.4% real rejections (~175 a month) and staying inside the carrier rule. Kill the block, keep the protection.

Why we cannot just do that: three locks on the door.

Lock	What holds the gate shut	What picks it
1 · Modifications	UW fixes our pre-filled data and applies the roof exclusion before approving: half of approvals change nothing, 26% get the RSE (85% of the time with no roof fact changing first), and six fields carry 90% of all attribute edits. Bind without the fixes and mispriced policies come back as NOCs at inspection.	This test: the RSE default plus agent-attested data on the six fields. Measured ceiling: 89% of approvals need zero UW edits.
2 · Rejections	The 8.4% of referrals rejected (~175 a month) are real - and 81% of the decline substance is not about age: commercial exposure (28%), historic districts (15%), unit counts (13%). No automated signal exists yet to catch any of it.	Pluribus signals: commercial in build now, historic 6-12 weeks, then shadow receipts against live referrals. Residual: ~19% of decline substance has no signal - ~30 quotes a month carriers must knowingly accept.
3 · Carriers	Pre-bind review of every 101+ dwelling is their requirement - even though ~2% of reviews find an age-driven problem. With the first two locks picked, removing the block is still their sign-off.	The receipts from locks 1 and 2, with the ~30/month residual priced up front.

Lock 1 has a menu; this test chooses from it. Context that matters: UW's "research" on these referrals is largely LandGlide and Zillow lookups - correct property data pulled by hand. Three ways to get that data onto the quote:

- **Buy it** - the parcel-data API we priced and passed on because it cost money. The same data UW pulls manually today.
- **Ask the agent** - the flip: six required fields, attested entry, near-zero build.
- **Keep the human** - the status quo. It works, and it is the OpEx and the attrition problem.

The test answers one question: **does the free option work, before we pay for the middle one?**

Be clear about day one: success changes nothing visible. The alert fires, the block stands, the referral is required. This is an evidence play - and the counterfactual shows why it matters. If Pluribus shipped tomorrow and this test had not run, we still could not stop blocking: the review would still own its data job and we would have no replacement for it. Success makes Pluribus the only remaining blocker and hands the carrier conversation its receipts - and it does not kill the alert; rejections and carriers still gate. It means that when the signals arrive, there is exactly one key left to ask for.

1 · The premise: a catch-all alert, and a review that mostly corrects our own data

DWELLING_AGE_MANUAL_REVIEW fires on every dwelling 101 years or older. Carrier guidelines require pre-bind review, so it blocks 100% of flagged quotes - roughly 7,600 fires a month, of which ~5,200 (68%) carry no other review-requiring alert. Of those, 40% get referred; 60% die without ever reaching an underwriter.

The BUC-4962 census told us what the review actually does. When agents ask, UW says yes 83% of the time and no 8.4%; 81% of decline substance is not about age (commercial exposure, historic districts, unit counts). On approvals, roughly half change nothing at all, and the most common substantive act - the roof-surfacing exclusion, 26% of approvals - is applied ~85% of the time without any roof fact changing first. It is a default, not a discovery.

The rest of UW's edits are corrections to data **we** put on the quote. We pre-fill dwelling attributes from vendors (Estated, 360Value, Smarty); on old homes the vendors are wrong; UW fixes our pre-fill after referral, with rate impact, behind frozen fields. The flip: stop pre-filling on 101+, make the agent enter the data, and UW stops being our data-entry QA.

2 · What the review's work actually is (warehouse run, Aug 7)

Universe: 10,620 DAMR-only, agent-referred, approved quotes (Jul 24 spine, terminal-action semantics); UW-attributed edits only. These are floors until the Q08 attribution bug closes - some "agent" edits were really UW.

The form is six fields, not twelve

Field UW edits on approval	Share of approvals
property_type	8.5%
roof_type (one field, stored under two keys)	~8.3%
total_sq_feet	8.2%
construction_type	8.1%
roof_shape	3.4%
year_built	3.0%
Everything else (stories, heating, bedrooms, baths, garages, foundation...)	0.3% or less - most never appear

Twelve fields were candidates from the code investigation; six carry the freight: 90% of every attribute edit UW makes (5,064 of 5,615), and 86% of attribute-edited quotes touch nothing outside them. The tail adds agent friction without adding coverage.

The ceiling: 89% of approvals need zero UW edits under the flip

What UW's edits were, per approved quote	Quotes	Share
Nothing outside {six-field form + RSE + derived noise} - the flip covers it	8,288	78.0%
Limit edits only (dwelling / vandalism / ordinance-law...) - 98.8% of these co-occur with an attribute fix or RCE recompute, i.e. repricing fallout, not judgment	1,204	11.3%
Deliberate coverage work - private-structures exclusions, trampoline, deductibles, address fixes	1,128	10.6%

Fix the attributes upfront and the reprice happens at quote time, so the limit "edits" never exist: 78.0% + 11.2% = **an 89.2% ceiling**. The last ~11% is real underwriting tailoring that better data does not remove - which is fine. The claim was never "no underwriter ever"; it is "no underwriter for data reasons."

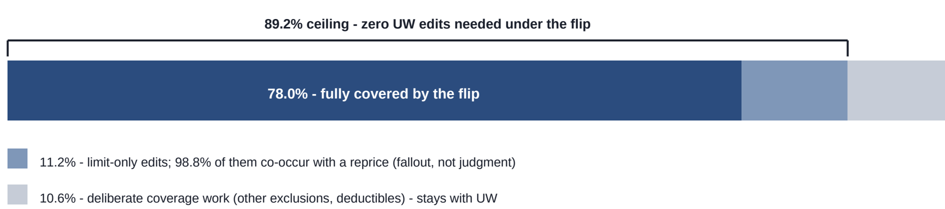

Inside the substantive-no-RSE bucket, the actor split confirms the shape. The six-field corrections run 85-95% UW-attributed - the flip's target work, on quotes where no exclusion was applied. The limit recomputes riding along are ~80% UW - repricing fallout. Coverage tuning (deductibles, personal-property and theft/water-backup limits) is majority agent - sale adjustments made during review, not review output. The deliberate UW residue is other exclusions and address fixes.

The premise now has a per-field receipt. Reconstructing each quote's per-version history (value at creation, value at referral, UW's review edits) across the 10,428 analyzable approvals: five of the six fields arrive pre-filled on 99-100% of quotes (roof shape is the exception at ~48% - it was never really a pre-filled field). Agents already modify them before referring on 9-23% of quotes per field - the flip makes today's optional entry mandatory and attested; it does not introduce agent entry. And when UW edits a field, 65-79% of that volume lands on pre-fill the agent never touched. "The review fixes our data" holds numerically.

One caution, flagged rather than papered over. Condition on the agent having touched a field, and UW still corrected it at 1.4-5.7x the rate of untouched pre-fill - sharpest on square footage, the #1 corrected field (25.3% vs 4.5%), where half of UW's fixes are re-fixes of something an agent just typed. This is observational and selection-biased: agents edit the quotes that are visibly broken, and today's voluntary edits are not tomorrow's attested required entry - so it cannot settle the flip either way. It does mean "agent enters it, UW stops fixing it" is unproven, and it is the data behind the county-validation backstop and the weekly garbage-in check.

3 · The test

	Treatment arm (101+ quotes)	Control arm
Roof-surfacing exclusion	Defaults on at quote creation; skipped when the condition score is rated Excellent; removable by referral with evidence	Today's flow - UW applies it during review
The six fields	Blank and required; the agent enters them	Pre-filled from vendors as today
Alert / block / referral / UW review	Unchanged - everything still refers, UW still reviews everything	Unchanged

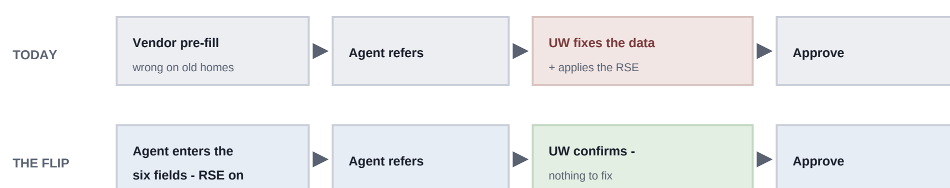

Same alert, same block, same referral - the only change is who supplies the data.

Run ~one month on a traffic split. Volume supports it: ~5,200 DAMR-only quotes a month. An optional third arm (flip without the RSE default) isolates which lever causes any damage. Build cost is small: the forced-entry mechanism already exists in code (requiresUserInput fires today on five fields when the pre-fill source is UNKNOWN) - the change is widening the predicate to "dwelling is 101+" and the key set to the six fields, plus the coverage-defaulting change for the RSE. Revert is a config flip. Measurement needs one eng ask: agent entries stamp source: USER in the app, but the per-field stamps are not mirrored to the warehouse today - wiring them through is a launch prerequisite for the test's per-field reads. Second open item: the RSE is a permission-gated coverage, so who can remove it needs an owner before launch.

4 · What it validates

One hypothesis: UW's modification work is caused by the pre-fill, not by judgment. The primary read is UW's edit rate on treated referrals - untouched approvals should climb from ~45-50% toward the 89% ceiling. If they do not move, the pre-fill was never the cause, whatever bind rate does.

Metric	What it guards against	Reads
Untouched-approval share (primary)	The hypothesis itself	In-window
Referral rate vs. the 40% baseline, split by reason	Confusion. RSE-removal referrals are by design; anything else rising is the signal	In-window
Form completion at quote start	Drop-off. Required fields land exactly where 60% of these quotes already die	In-window
Bind rate	Net business impact	In-window
Agent-entered value distributions vs. vendor/county priors	Garbage-in: plausible junk typed to clear the form	Weekly
NOC (notice of cancellation) rate on the tagged cohort	Mispriced binds surfacing at inspection	A quarter out

5 · If it succeeds

Treated referrals come back near the ceiling and the guardrails hold. The review's data job is gone: an old home with agent-attested attributes and the exclusion defaulted does not need a human on the approve path. What follows:

- **The carrier case gets its receipts.** "Half your required reviews changed nothing before; now ~89% change nothing, and the remaining 11% is named coverage tailoring." That is the ask to drop or condense pre-bind review on the clean path.
- **Conversion proceeds on the adopted sequence.** The modifications leg is proven; the alert still cannot die until the rejection sub-alerts have shadow receipts. The test does not touch that leg.
- **County records demote.** The agent is the source of record; study #3 becomes an optional backstop or a post-bind check, and Rob's build-vs-buy comes off the critical path.

A half-success - edit rate moves but drop-off spikes - means the hypothesis is true and the form is too expensive. Soften the form (fewer required fields, smarter defaults); do not abandon the idea.

6 · If it fails

UW keeps editing agent-entered data at the old rate. The corrections were never about vendor pre-fill: UW's desktop research beats what agents know or bother to type. What follows:

- **The agent cannot be the source of record.** The source-of-record gap is real and external. County/assessor matching returns as the load-bearing mechanism, and Rob's build-vs-buy is the critical path again.
- **The modifications leg of conversion stays blocked.** What survives from this design is the RSE default and the rejection sub-alerts.
- **Failure is granular, not binary.** Per-field USER stamps say which flips worked - once the stamp mirror ships (the stamps live in the app today, not the warehouse). If year built and square feet come back clean while construction type still gets corrected, keep two and revert one.
- **The ugly failure is garbage-in:** agents typing plausible junk to clear the form - worse than today, because mispriced quotes bind and surface as NOCs at inspection - and the observational data already leans here: UW corrects agent-entered values at 1.4-5.7x the untouched rate (selection-biased, but a warning). This is why the value-distribution check runs weekly, and why this failure mode argues for county *validation*, not pure attestation.

The asymmetry is why the test is worth running: it is a config-flip build and both outcomes pay. Success buys the carrier conversation. Failure proves the county-data investment is necessary - a fact you would otherwise spend a quarter and a vendor contract discovering.